yql

techtalk

data processing at scale

20120927 – mirek@yahoo-inc.com

the plan

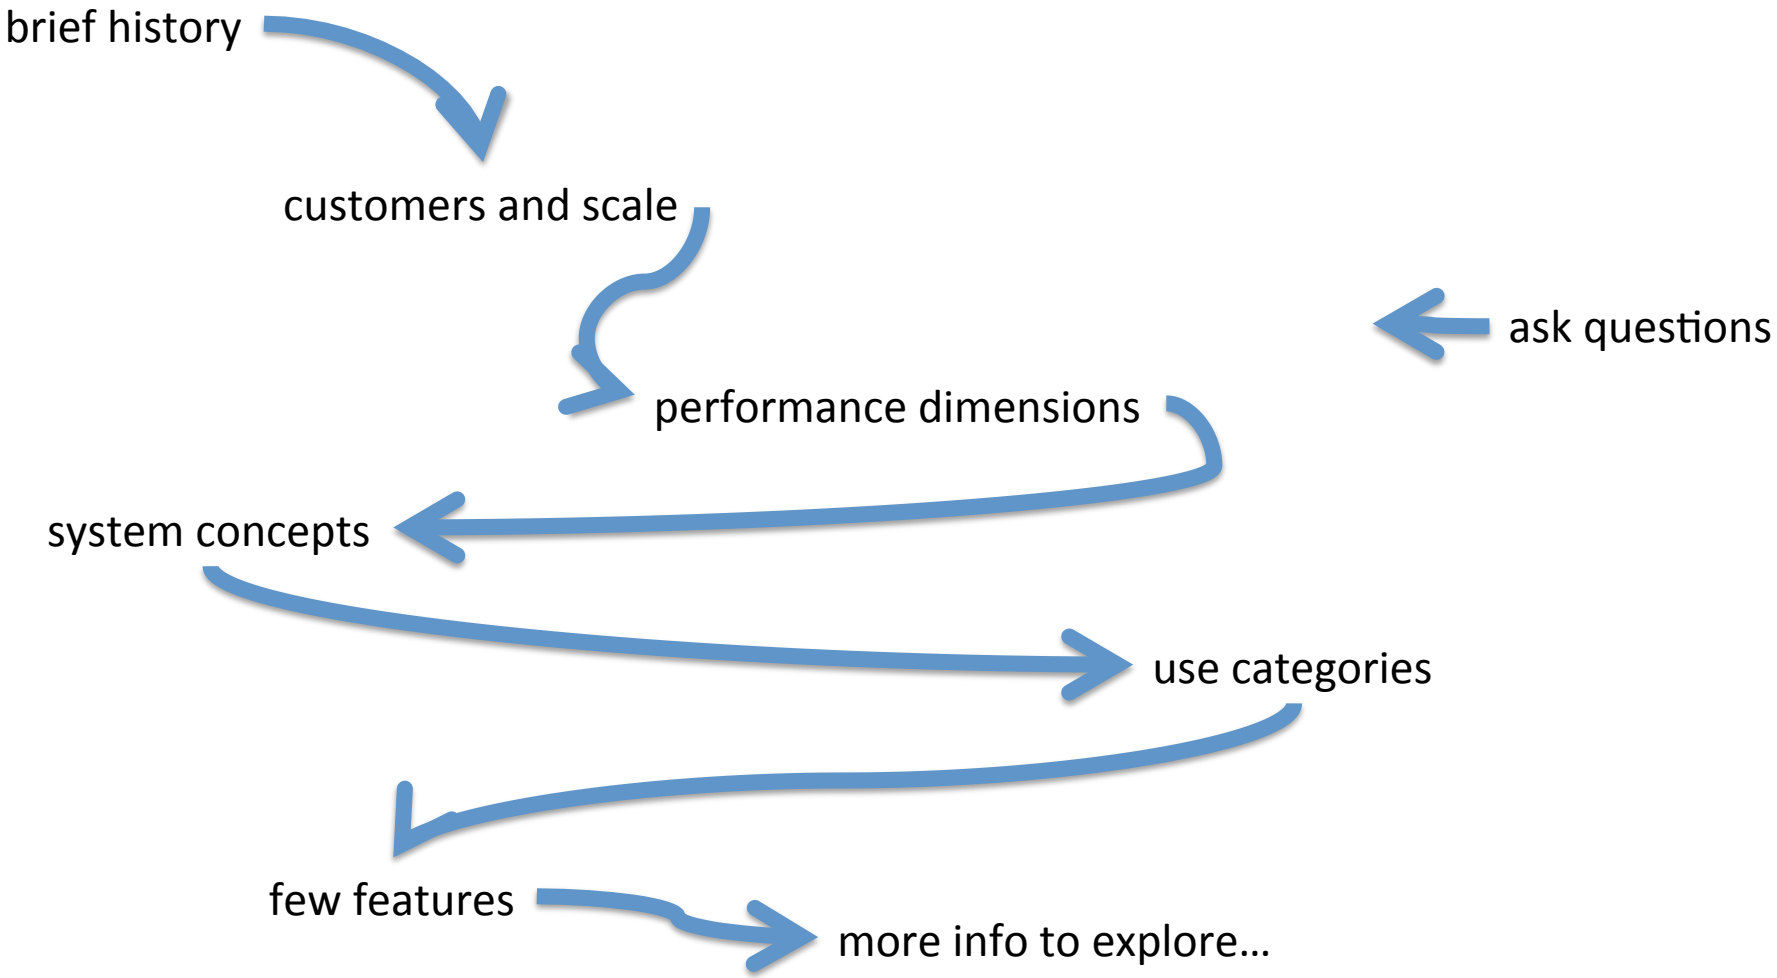
brief history
customers and scale
ask questions
performance dimensions
system concepts
use categories
few features
more info to explore…

brief history

yql started as an upgrade to  (developers asked for cmd version of pipes)

evolved into “select * from internet” – cloud data serving

launched October 2008

added insert update delete – grew its vocabulary

<execute> added

developer console

…lots of features

served 5B queries September 2010

upgraded runtime architecture to scale

introduced tenancy into serving

over 1200 open tables

served 20B queries September 2011

over million of pipes

hundreds of yahoo! tables added

over hundred of on-boarded customers

today

the mission: make data discovery, enrichment and serving fast and easy

customers

three categories of customers

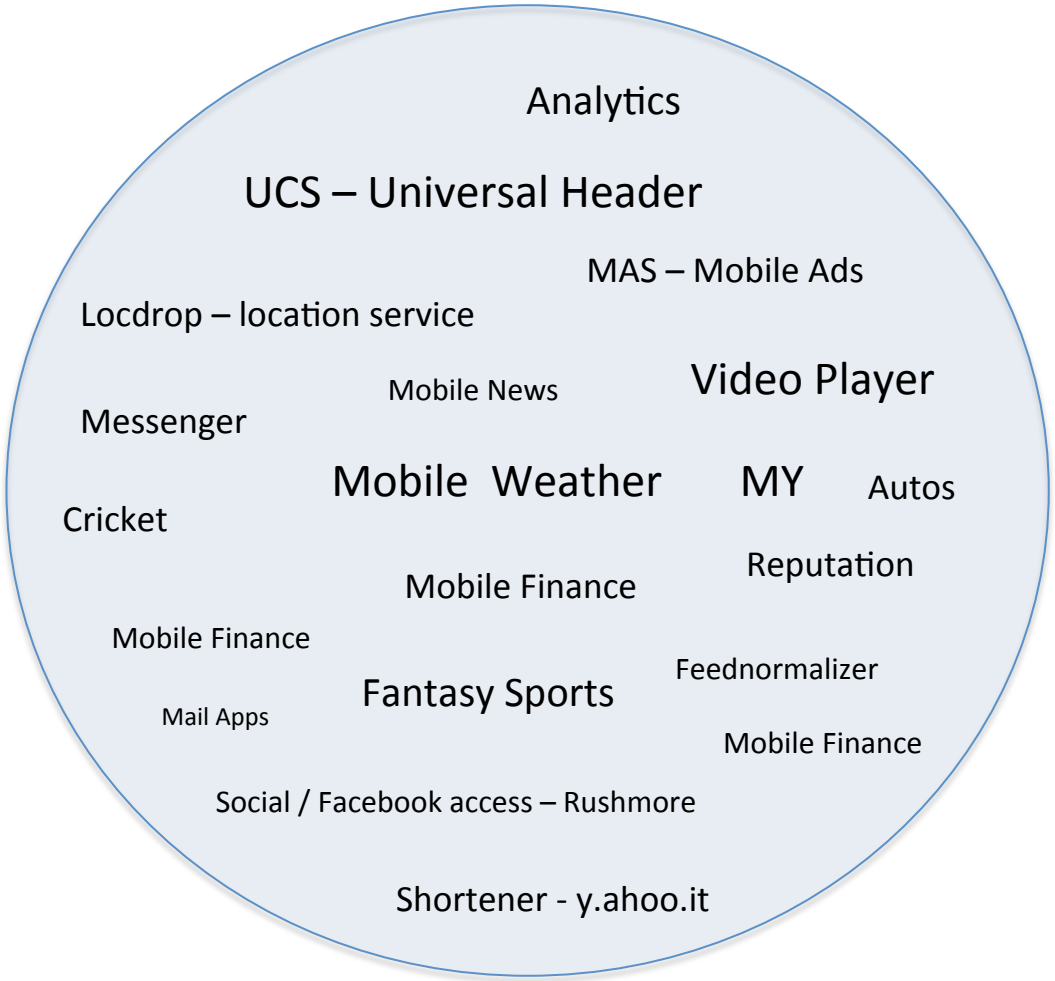

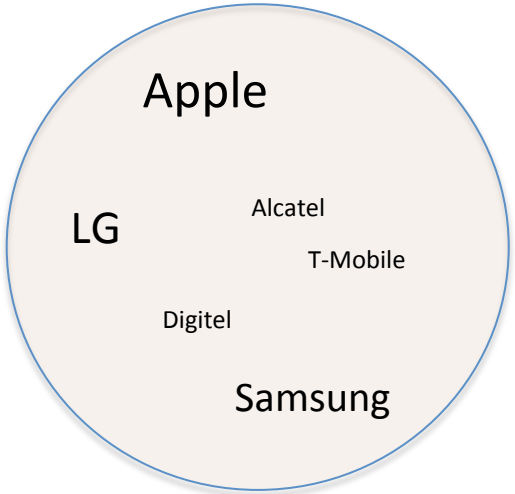

yql query volume in 2012

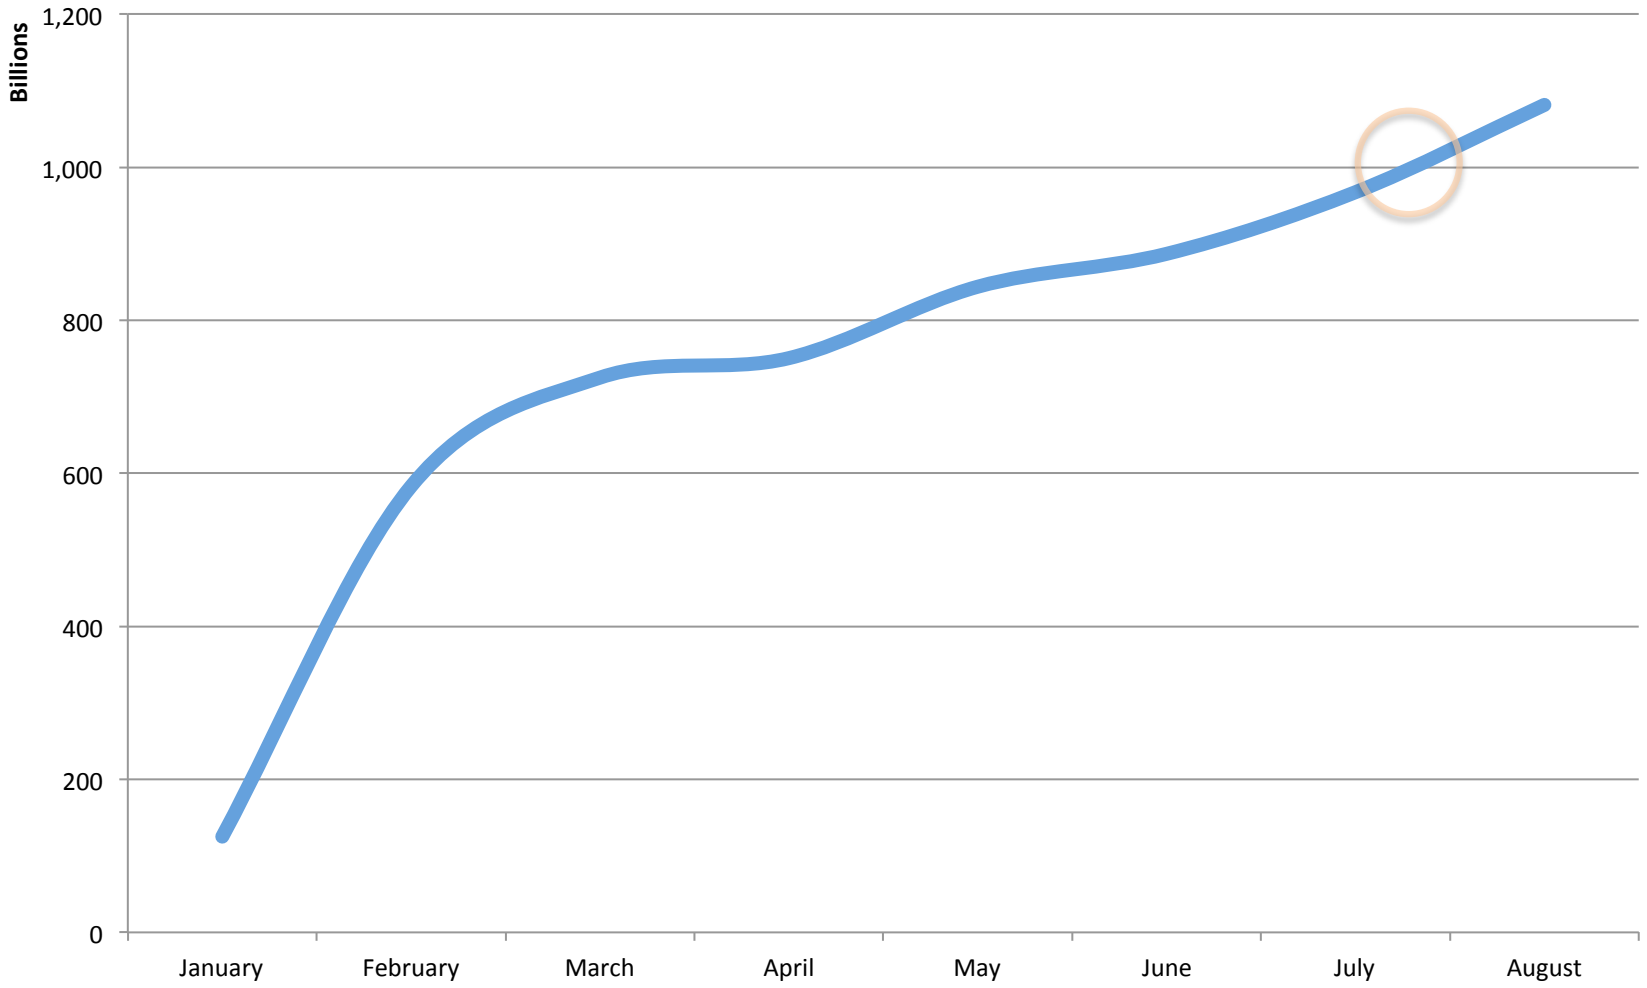

Apple

major partner – grows with iOS adoption

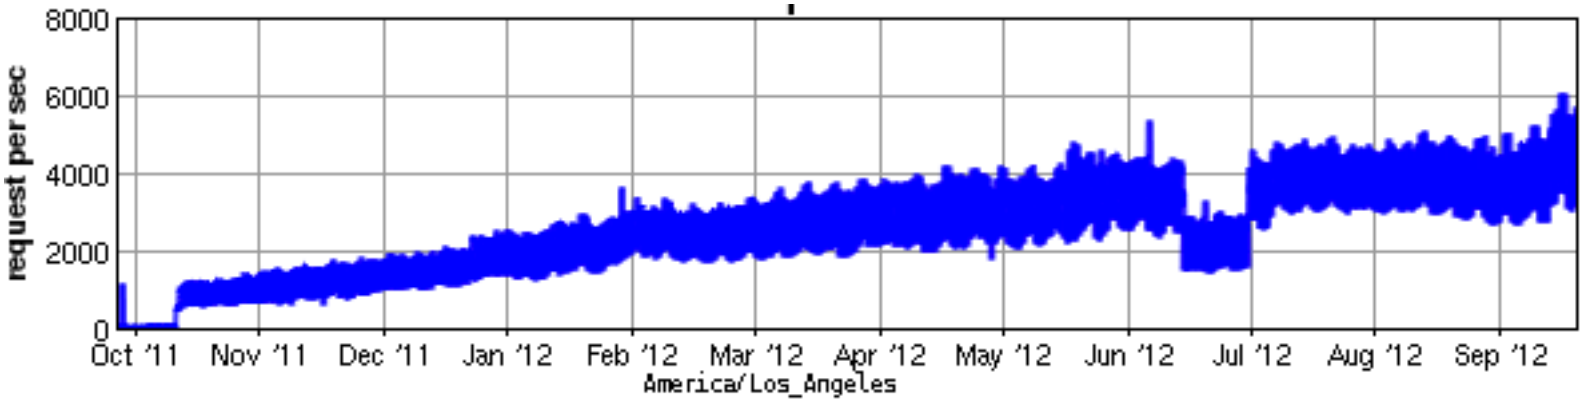

predominantly weather / siri

heavily utilizes yql.multi

Apple

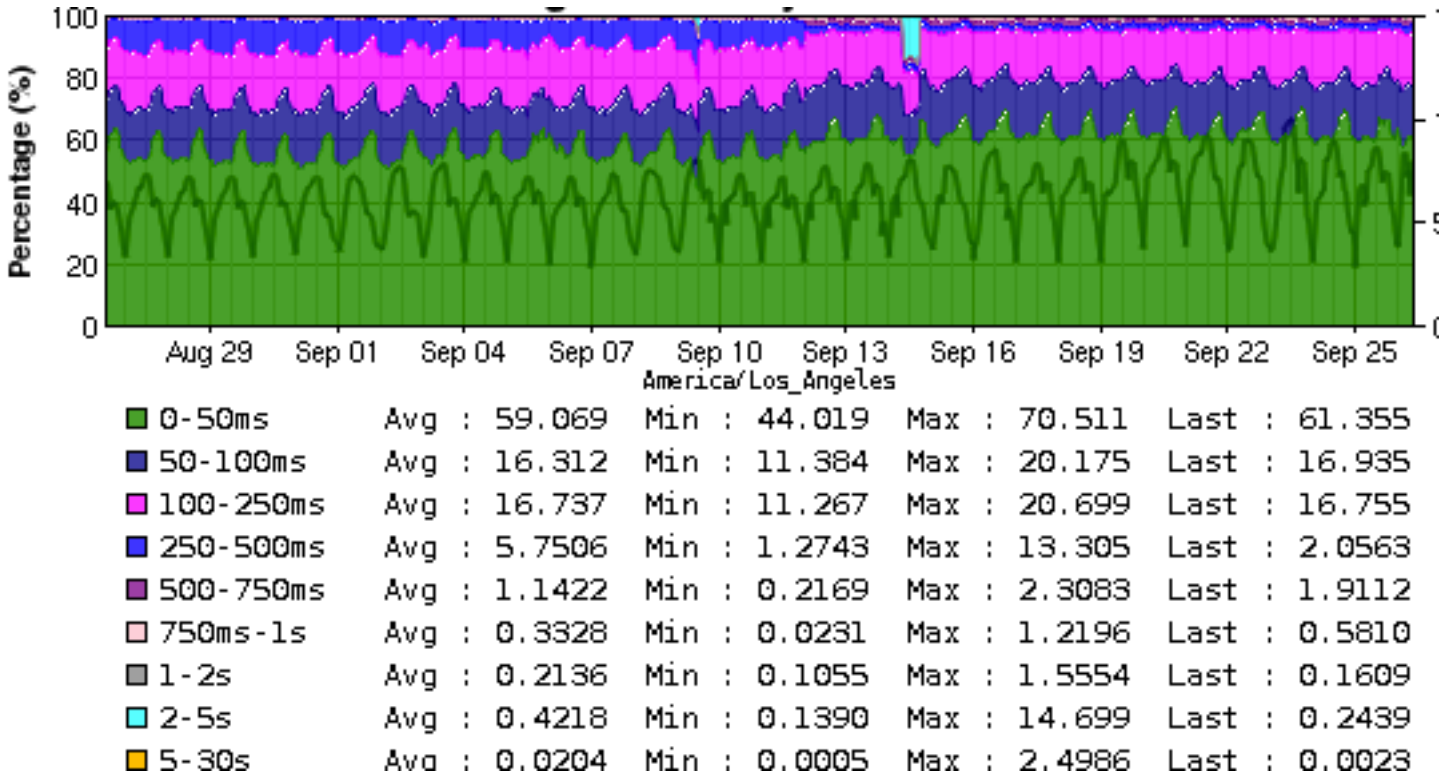

places performance demands on the platform

UCS – Universal Header

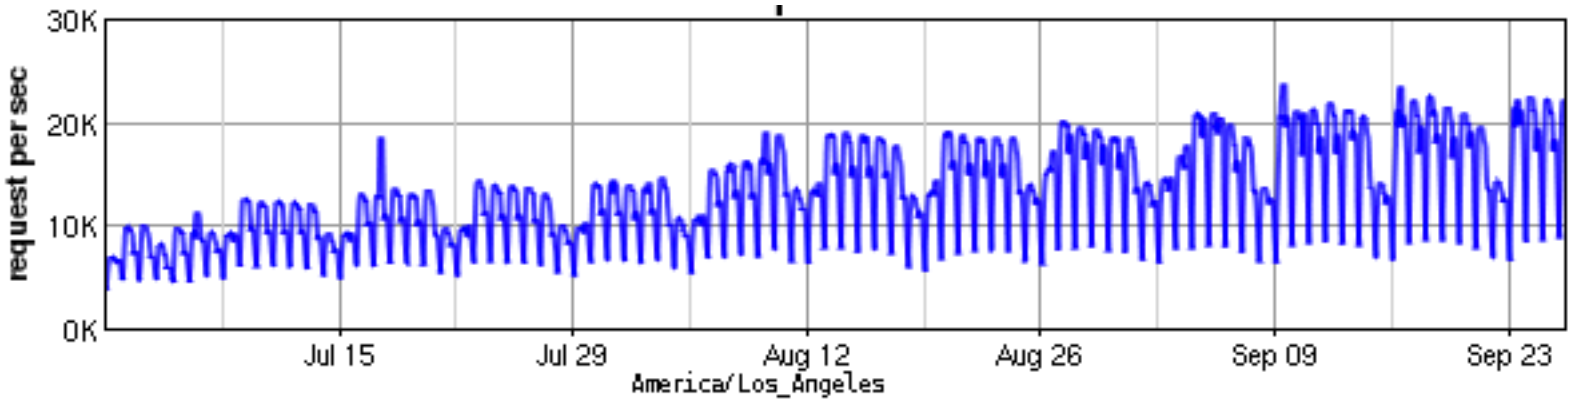

part of nearly every yahoo page

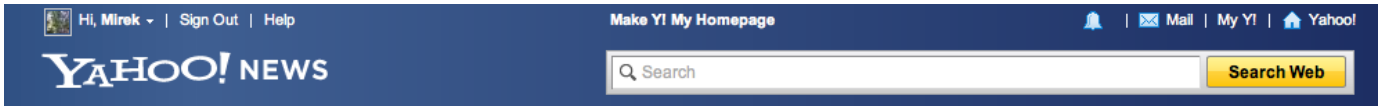

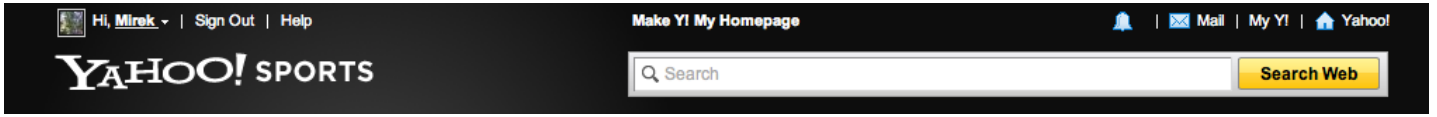

UCS – Universal Header

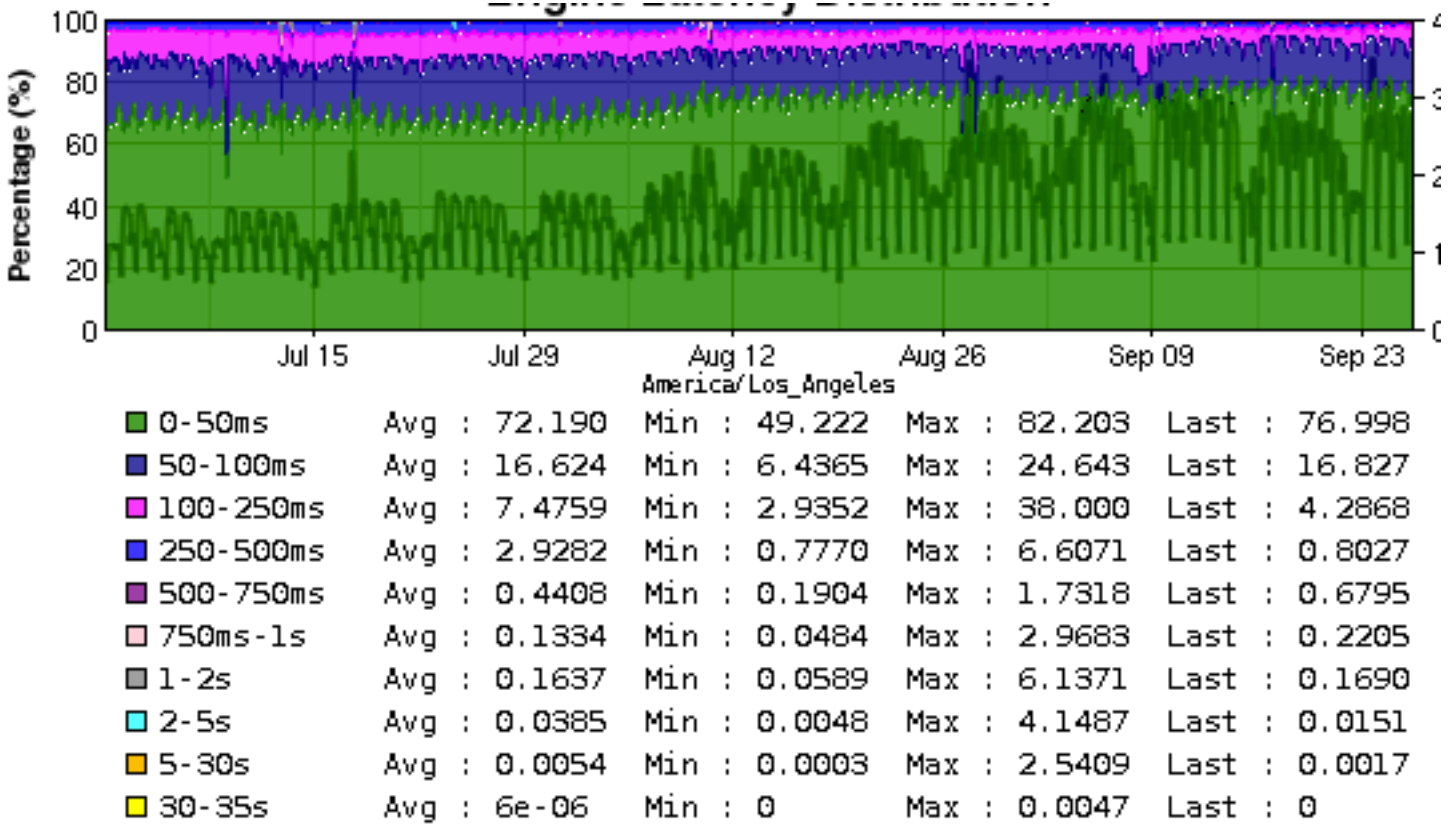

performance

lets look at platform performance as function of three dimensions

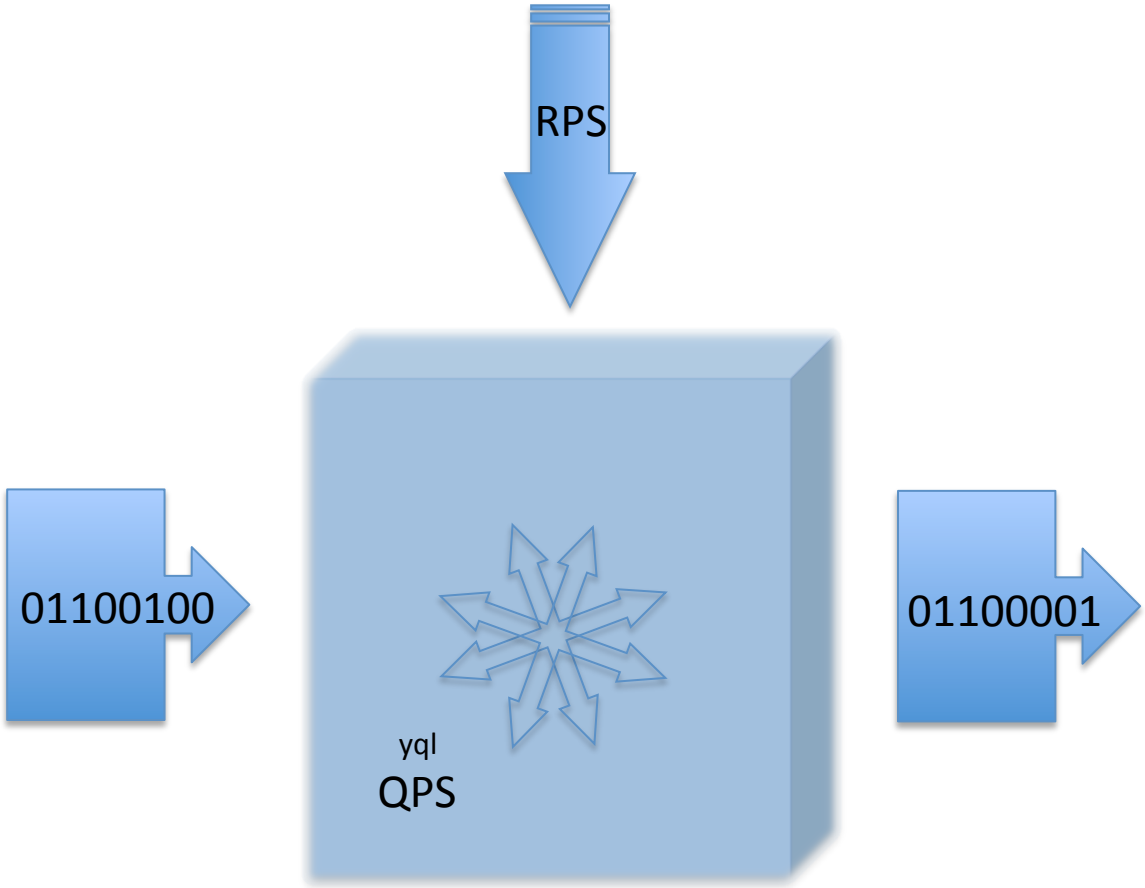

requests per second

represents inbound demand for work and product

efficient system does work fast and cheap

~35K average

~60K at peak

few* platforms at y! do this much work over HTTP

what's in a yql query?

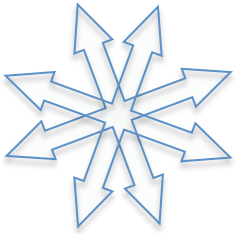

compile statement

gather response data

$$q == pr + cs + bep + grd + ppr$$

parse request

build execution pipeline

package product and return

over 400K / sec

yql queries represent work done by the engine

1T / month

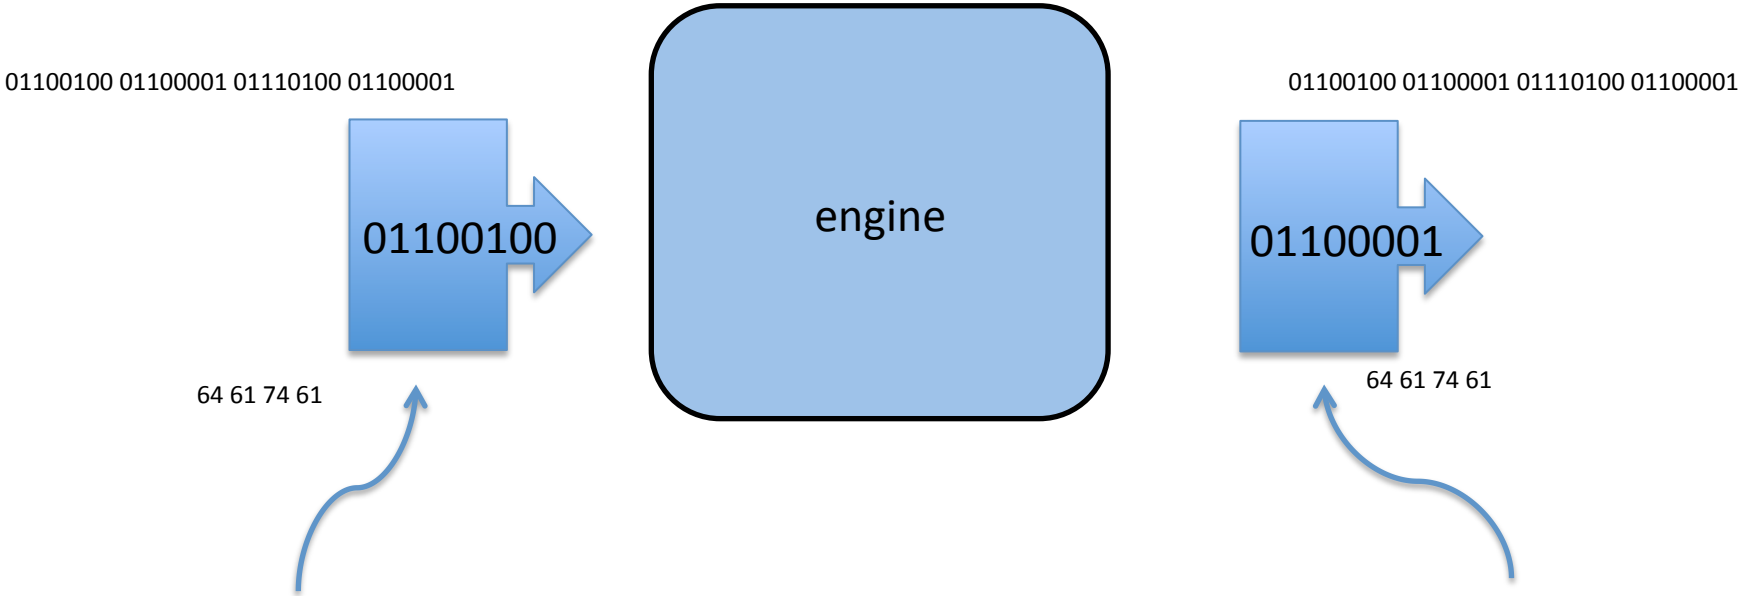

inbound over 30TB / day

produces over 10TB / day

or

about 100M really long articles

data volume processed by the engines

data flow by platform component per day

45TB in

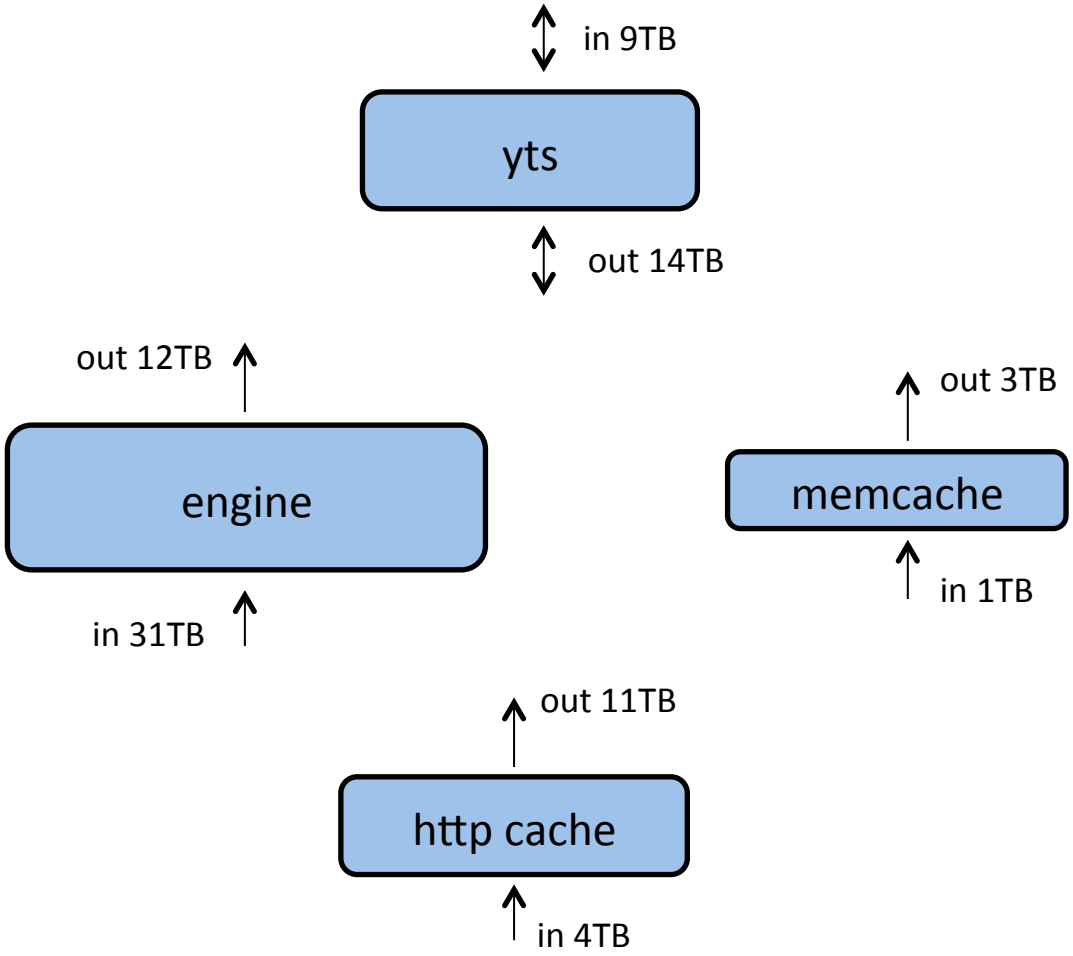

40TB out

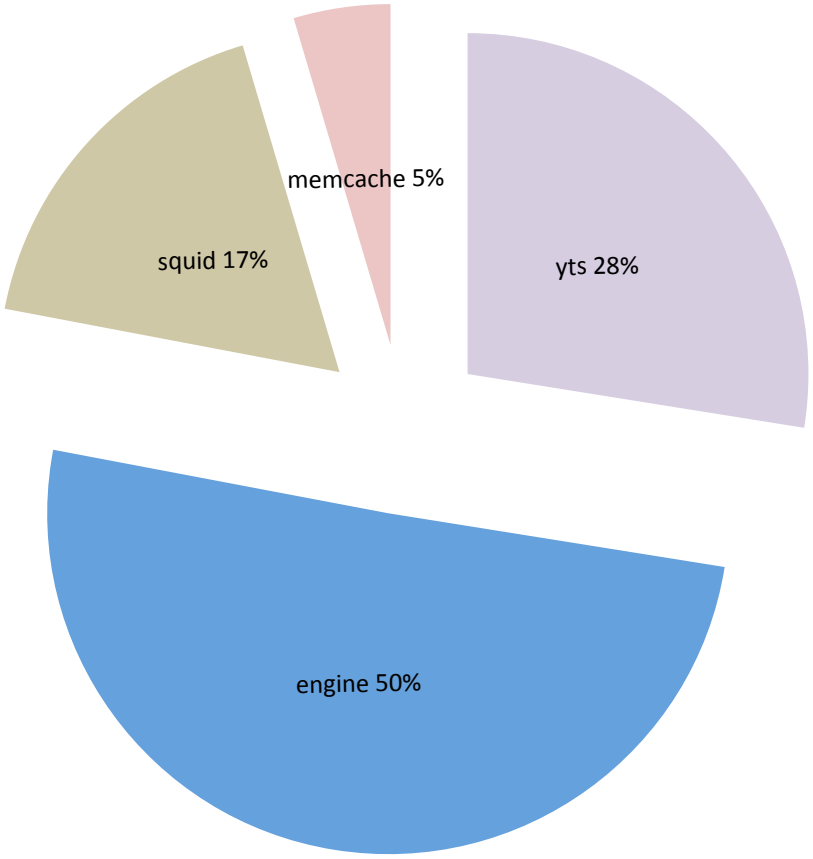

data seen by the platform component

yql platform sees over 2.5 PT in a month

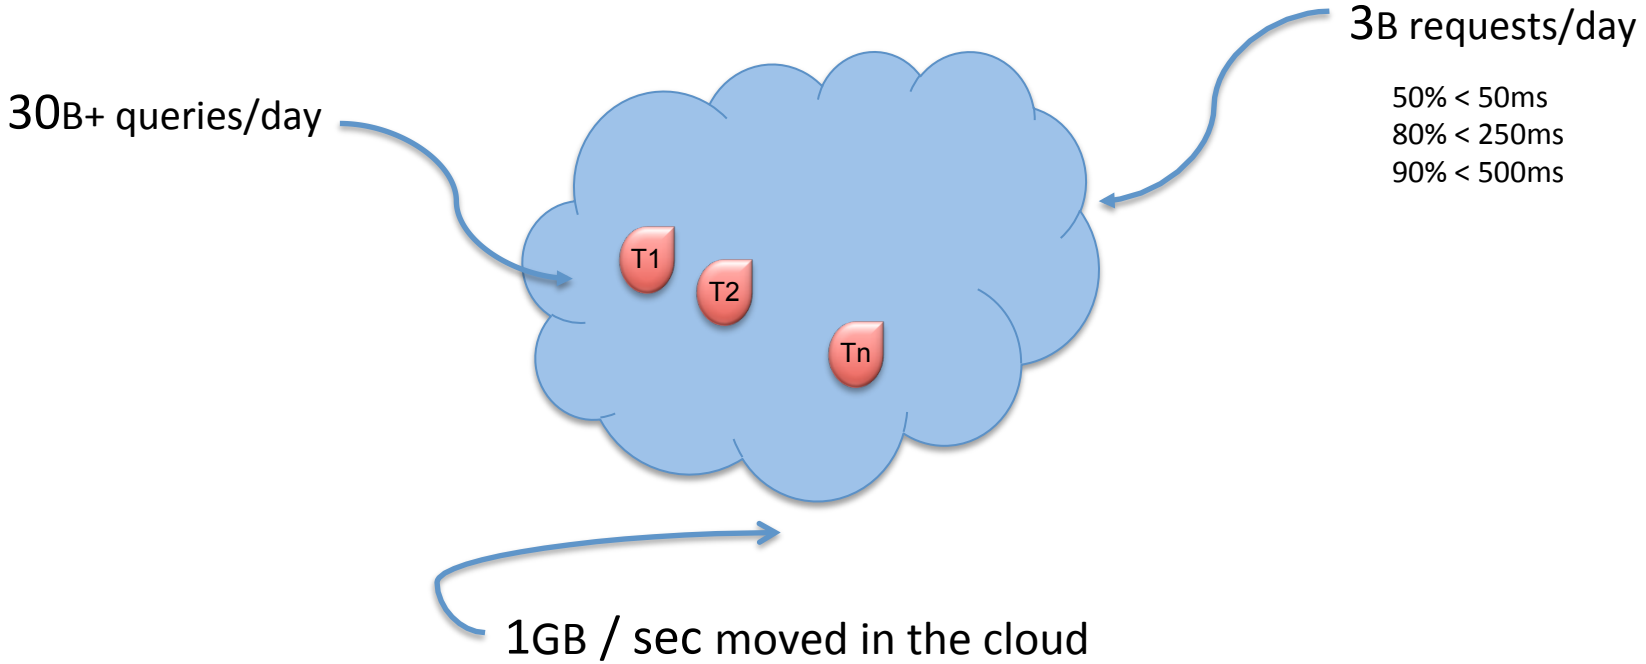

yql cloud = all yql system components

system

- Deployment
  - Active in 8 colos globally
    - US: SP2, MUD, AC4, BF1 (GQ1, NE1) = 70%
    - EU: CH1, UKL, IRD = 15%
    - AP: SG3 = 15%

- Scales
  - horizontally
  - vertically

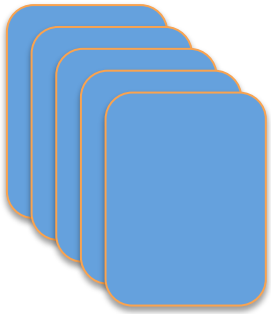

yql platform design reflects data supply and demand model

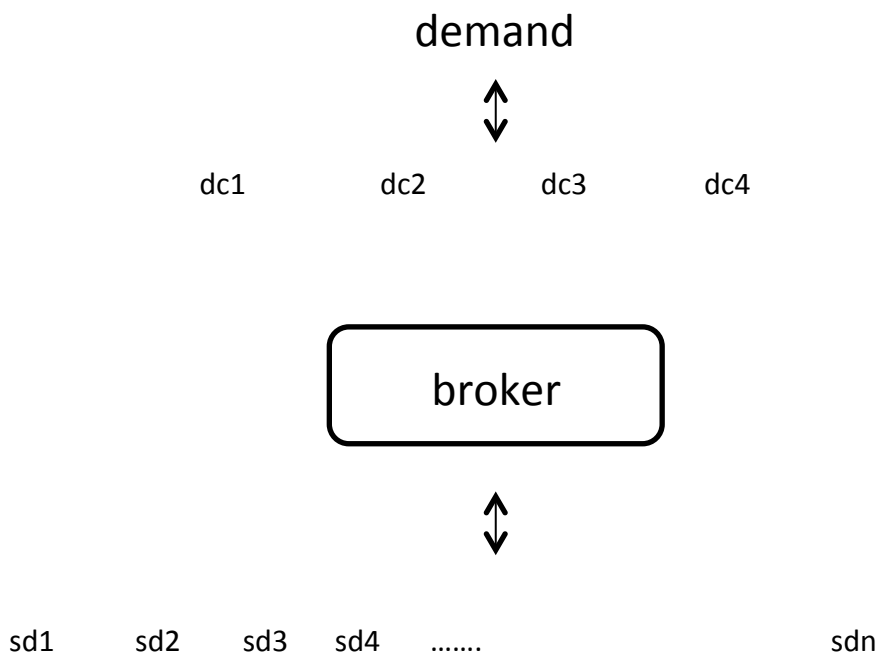

what does the production system look like

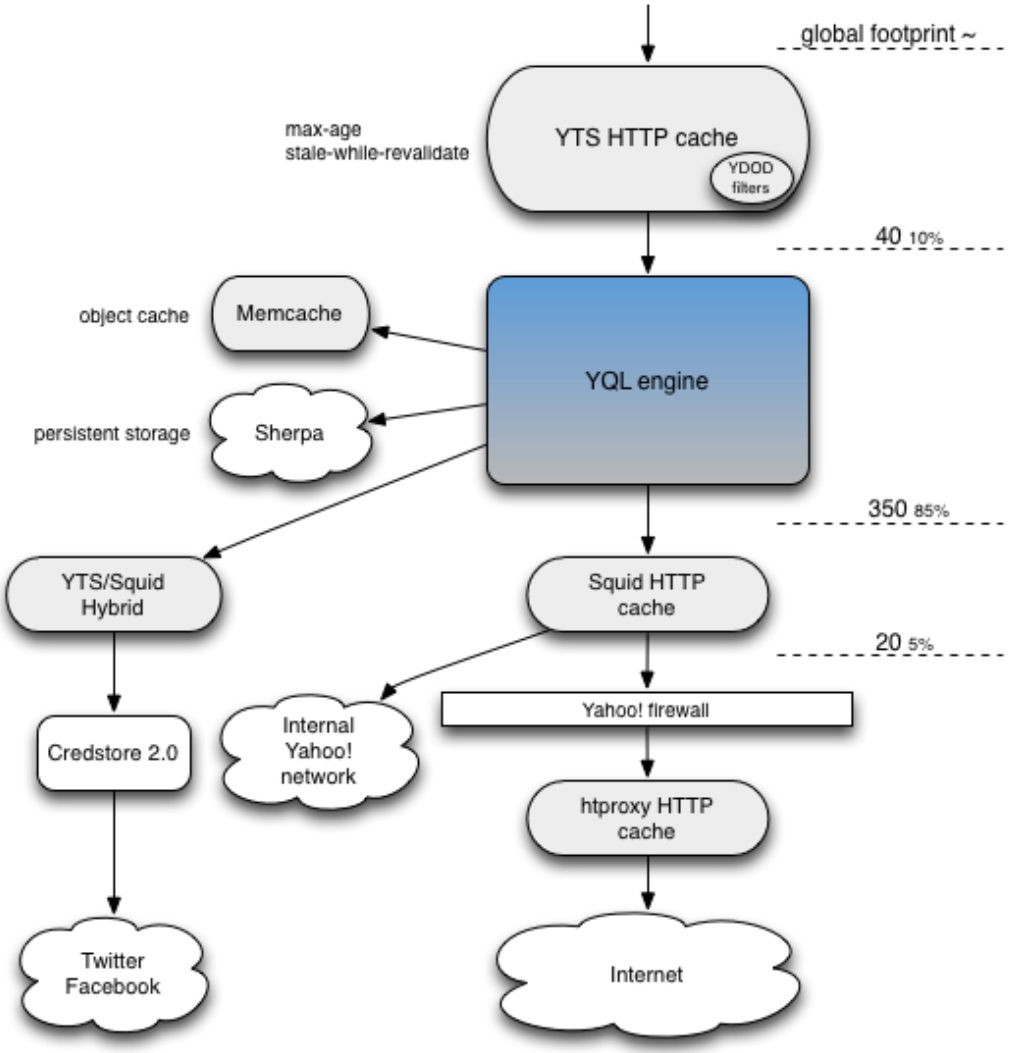

yql engine request view

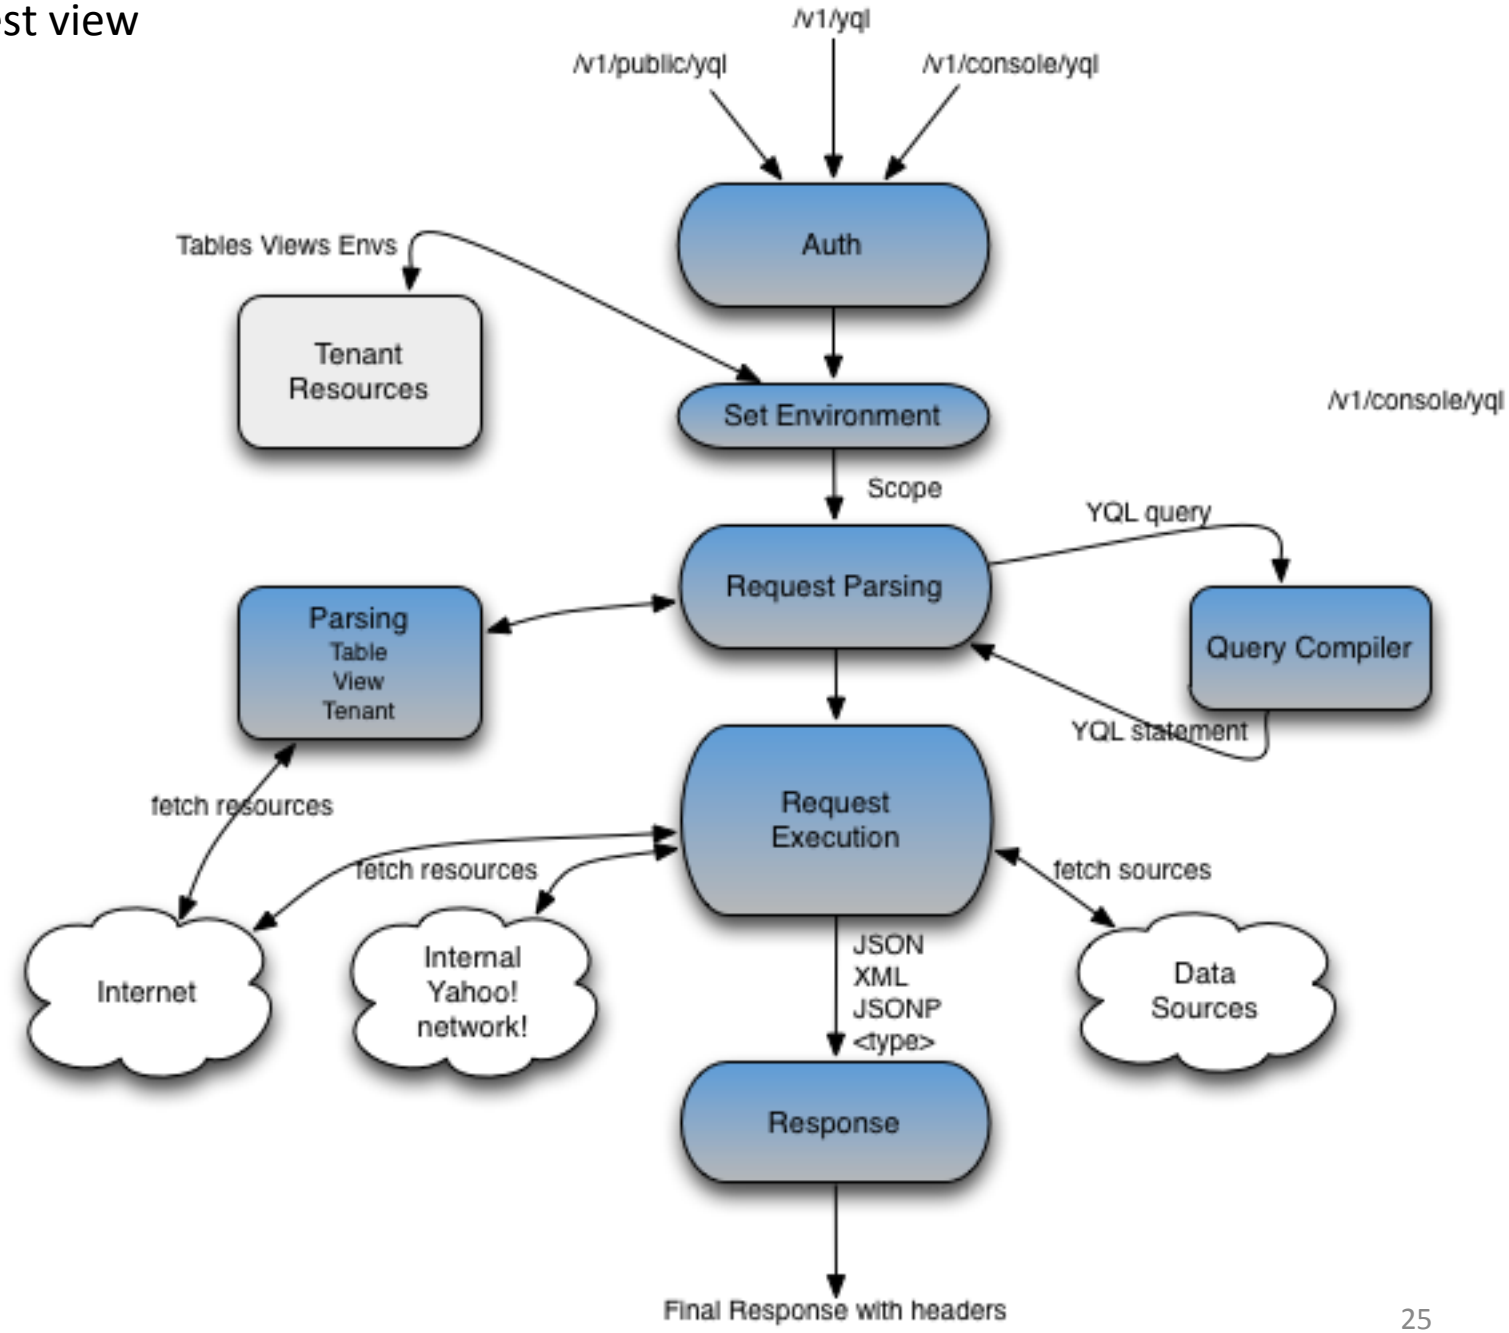

yql container view

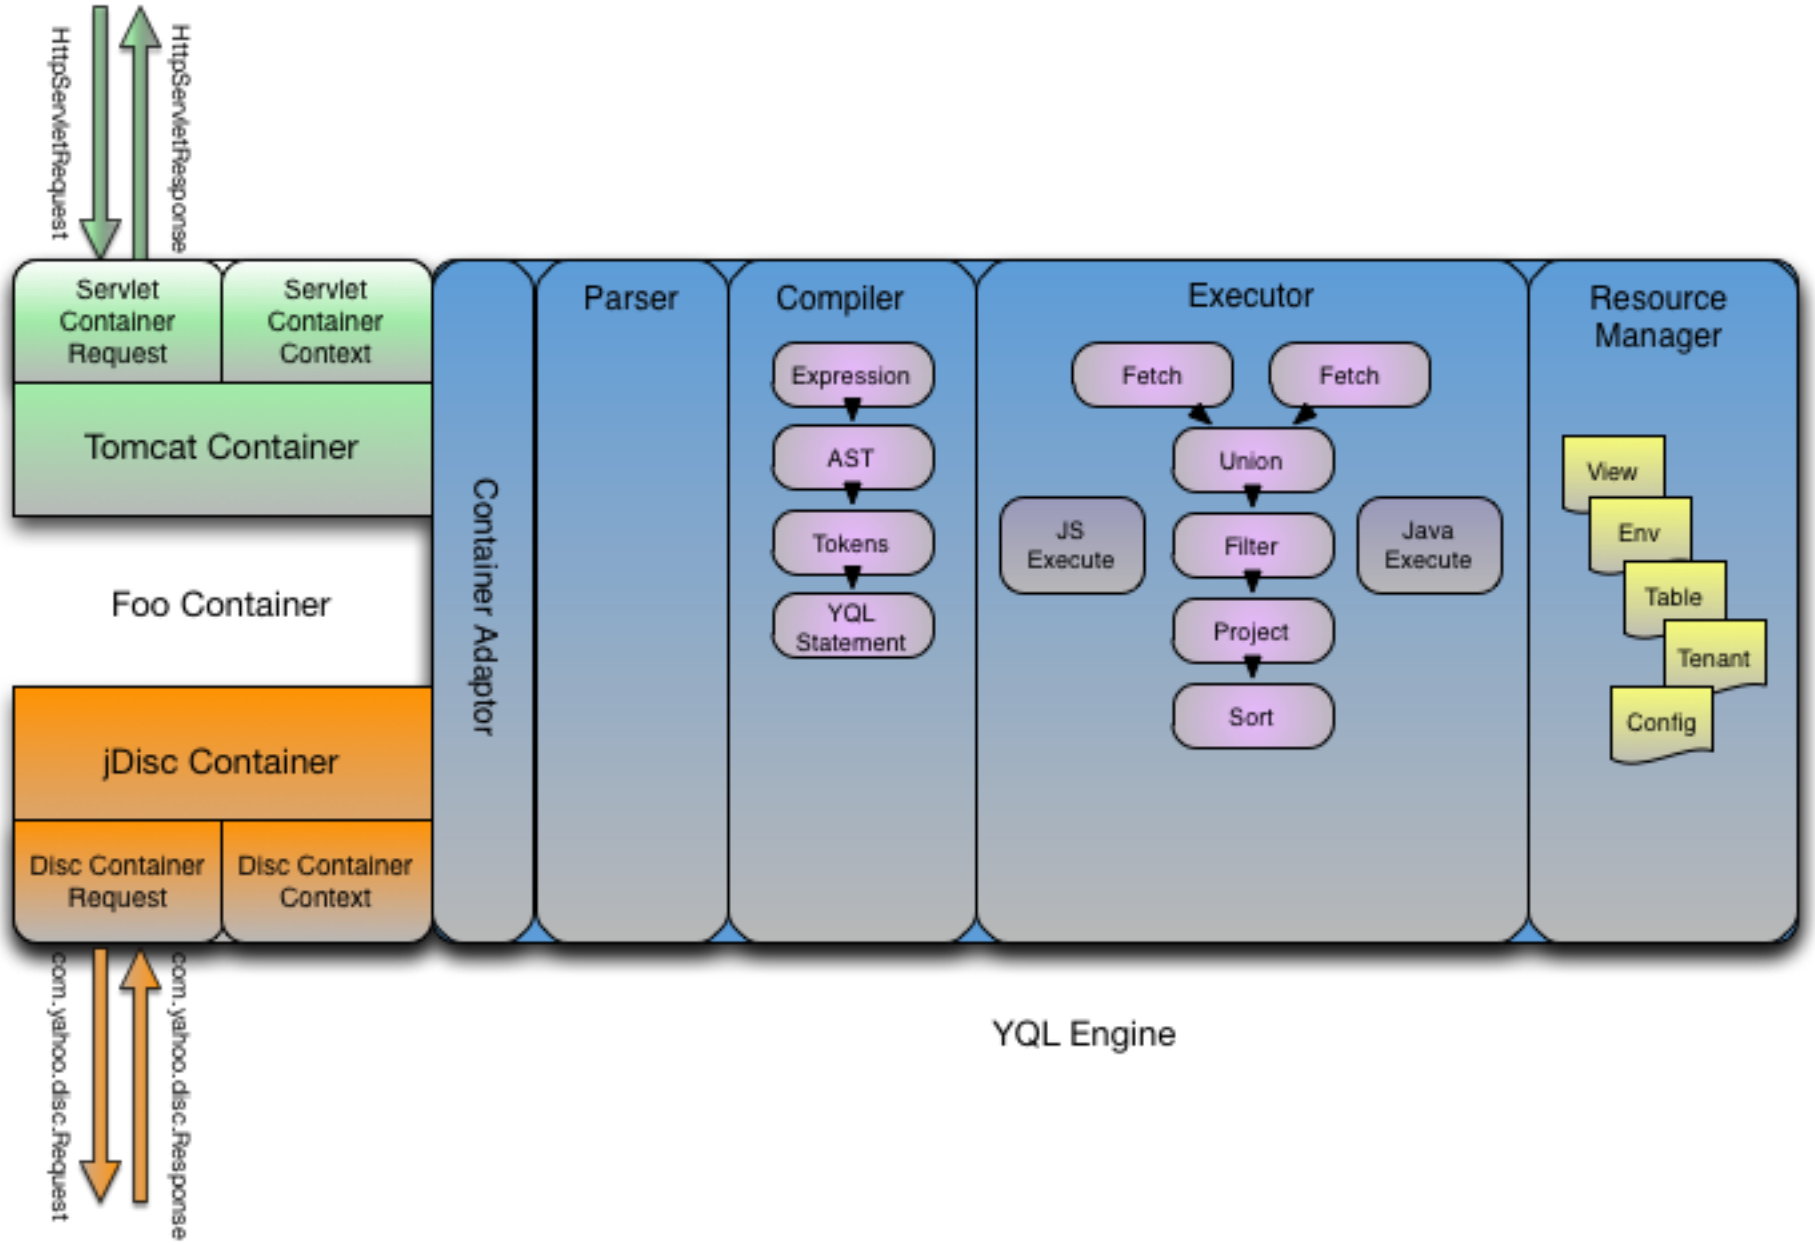

uses

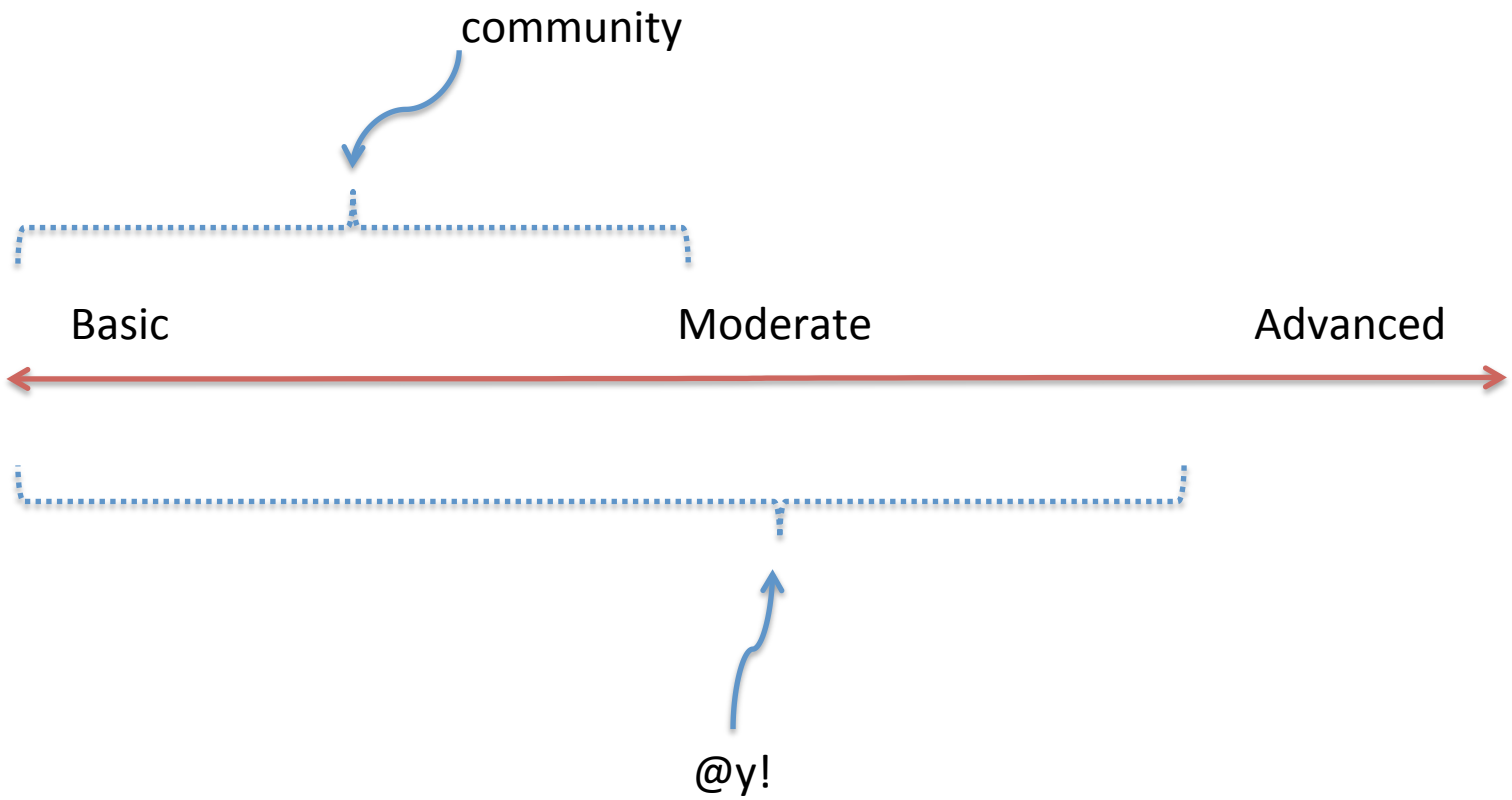
community
Basic
Moderate
Advanced
@y!

Simple – access data source and use syntax features

```
select * from rss where url
    in (select title from atom where url="http://y.ahoo.it/lu5F/")
        and description like "%Wall Street%" limit 10 | unique(field="title")
```

Moderate – use execute to shape data

```
delete from twitter.status
    where id="2108869549" and username=@username and password=@password
```

```
<delete itemPath="" produces="XML">
    <urls>
        <url>http://twitter.com/statuses/destroy/{id}.xml</url>
    </urls>
    <inputs>
        <key id="username" type="xs:string" required="true" paramType="variable"/>
        <key id="password" type="xs:string" required="true" paramType="variable"/>
        <key id="id" type="xs:string" required="true" paramType="path"/>
    </inputs>
    <execute><![CDATA[
        y.include("http://yqlblog.net/samples/base64.js");
        var authheader = "Basic "+Base64.encode(username+":"+password);
        response.object = request.header("Authorization",authheader).del().response;
    ]]></execute>
</delete>
```

Advanced – native implementation of data API

```
insert into yql.storage.admin (url) values ("http://localhost/table.xml")
```

not a use

caching only

(if you use yql for its cache then something went seriously wrong elsewhere in your design and implementation)

features

there are many features

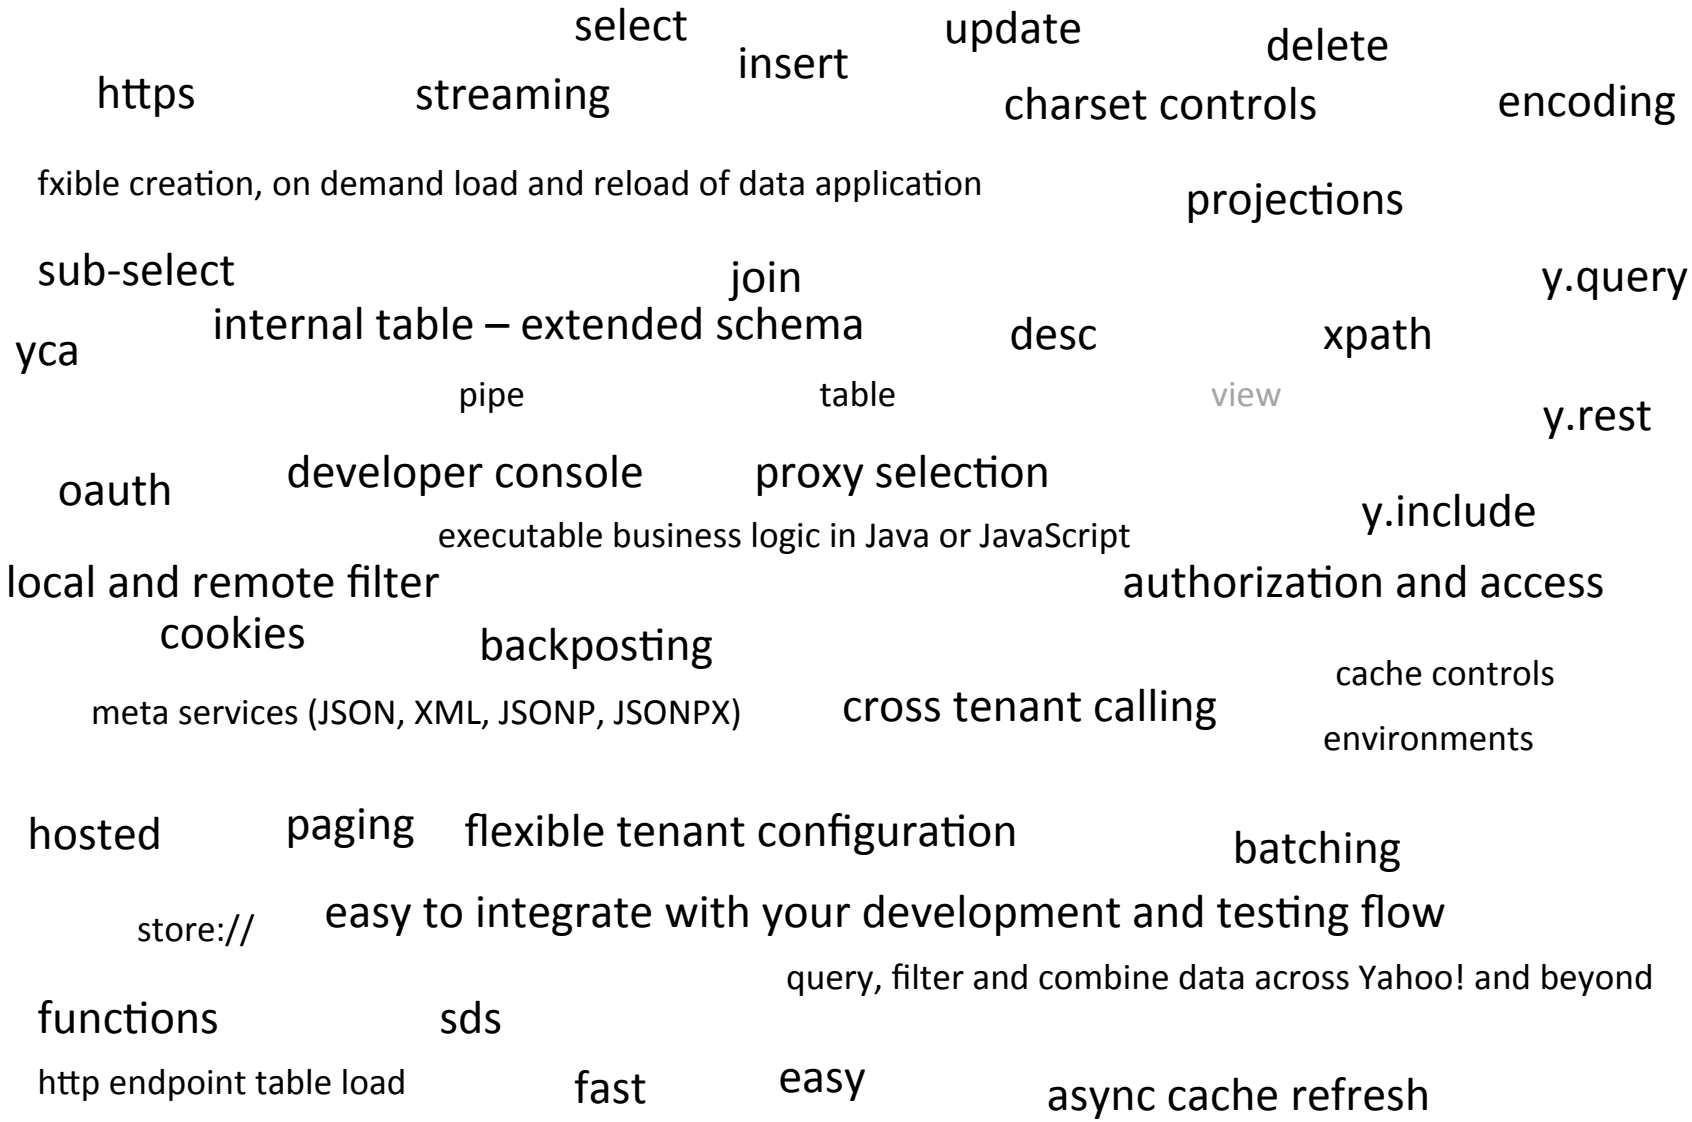

master those

pipe table view

pipe

a workflow

encapsulation of multiple tables and queries

expressed through JSON definition

created through GUI

instruction set bundle with operators for yql engine

available from execute y.pipe("id")

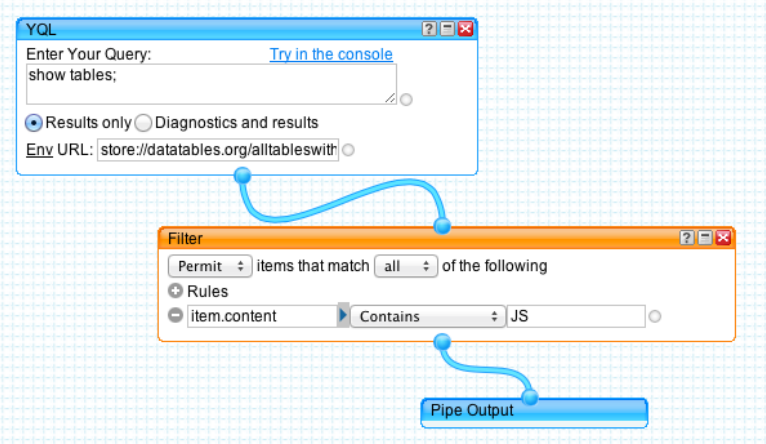

table

yql data app definition / plugin

runtime instruction container for yql

fundamental way to express source binding

cross tenant calling

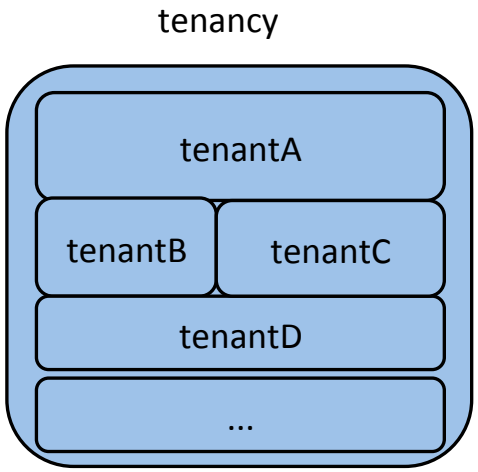

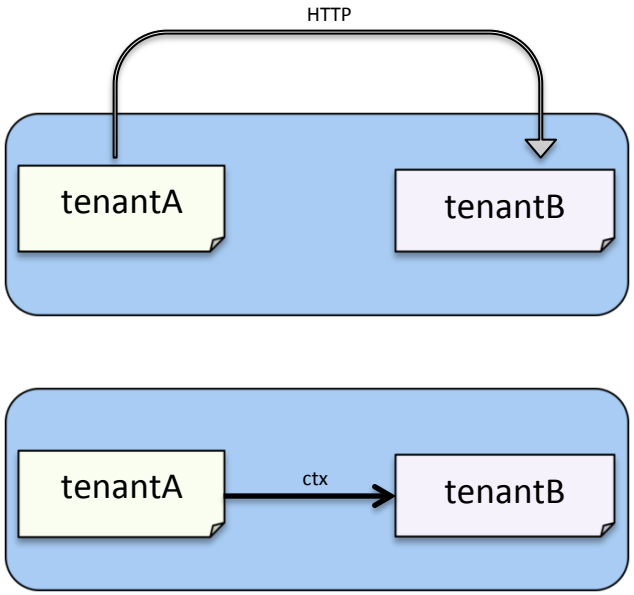

- keeps client resources isolated from each other
  - Tables, Views, Envs
  - Secrets (keydb)
  - Certificates (yca)
  - Customized OC TTL
- authorization schemes: public, oauth, cookie/crumb, YQL token
- custom hostnames: **media.query.yahoo.com**
- tables access can be restricted per tenant
- generate default request characteristics (b-cookie, user-agent, cache control)

coming soon…

views

tenant self-provisioning

even more flexible tenant operated resource reload

…..

dblink

remote table execute

recap

yql dna is ...

a pipes effort funded by Jerry and Filo in 2005

grew this investment slowly

multi tenant cloud data serving platform

serving over 1T queries/month
while processing over 2PT

more info...

on-boarding: easy

just search backyard for yql on-boarding

do not launch with public rate limited tenant – **query.yahooapis.com**

mirek@yahoo-inc.com

yql-discuss@yahoo-inc.com

devel.corp.yahoo.com – internal yql docs

know Java well...? yql could be a place to write it – let us know